



Course Name: Data Structures and Applications

Course Code: BCS304


Module 4
Trees

Binary Search Tree

Go, change the world

- Heap
 - a min (max) element is deleted. $O(\log_2 n)$
 - deletion of an arbitrary element $O(n)$
 - search for an arbitrary element $O(n)$
- Binary search tree
 - Every element has a unique key.
 - The keys in a nonempty **left subtree** (**right subtree**) are **smaller** (**larger**) than the key in the root of subtree.
 - The left and right subtrees are also binary search trees.

Examples of Binary Search Trees



Searching a Binary Search Tree

```
tree_pointer search(tree_pointer root,  
                    int key)  
{  
    /* return a pointer to the node that  
       contains key. If there is no such node,  
       return NULL */  
  
    if (!root) return NULL;  
    if (key == root->data) return root; if  
    (key < root->data)  
        return search(root->left_child,  
                       key);  
    return search(root->right_child, key);  
}
```


Another Searching Algorithm

```
tree_pointer search2 (tree_pointer tree, int
    key)
{
    while (tree) {
        if (key == tree->data) return tree; if
            (key < tree->data)
                tree = tree->left_child; else
                tree = tree->right_child;

    }
    return NULL;
}
```

$O(h)$

Insert Node in Binary Search Tree



Insert 80


Insert 35

Insertion into A Binary Search Tree

```
void insert_node(tree_pointer *node, int num)
{tree_pointer ptr,
    temp = modified_search(*node, num); if (temp ||
    !(*node)) {
    ptr = (tree_pointer) malloc(sizeof(node)); if (IS_FULL(ptr)) {
        fprintf(stderr, "The memory is full\n");
        exit(1);
    }
    ptr->data = num;
    ptr->left_child = ptr->right_child = NULL;
    if (*node)
        if (num<temp->data) temp->left_child=ptr; else temp-
            >right_child = ptr;
    else *node = ptr;
}
```


Deletion for A Binary Search Tree

Go, change the world




Deletion for A Binary Search Tree


non-leaf
node



Before deleting 60



After deleting 60




Definition

- A **graph** G consists of two sets
 - a finite, nonempty set of vertices $V(G)$
 - a finite, possible empty set of edges $E(G)$
 - $G(V,E)$ represents a graph
- An **undirected graph** is one in which the pair of vertices in a edge is unordered, $(v_0, v_1) = (v_1, v_0)$
- A **directed graph** is one in which each edge is a directed pair of vertices, $\langle v_0, v_1 \rangle \neq \langle v_1, v_0 \rangle$

tail $\xrightarrow{\hspace{1.5cm}}$ **head**

Examples for Graph

Go, change the world




G_1

complete graph

$$V(G_1) = \{0, 1, 2, 3\}$$

$$V(G_2) = \{0, 1, 2, 3, 4, 5, 6\}$$

$$V(G_3) = \{0, 1, 2\}$$




G_2

incomplete graph

$$E(G_1) = \{(0, 1), (0, 2), (0, 3), (1, 2), (1, 3), (2, 3)\}$$

$$E(G_2) = \{(0, 1), (0, 2), (1, 3), (1, 4), (2, 5), (2, 6)\}$$

$$E(G_3) = \{<0, 1>, <1, 0>, <1, 2>\}$$



G_3

complete undirected graph: $n(n-1)/2$ edges

complete directed graph: $n(n-1)$ edges



Complete Graph

- A complete graph is a graph that has the maximum number of edges
 - for **undirected graph** with n vertices, the maximum number of edges is $n(n-1)/2$
 - for **directed graph** with n vertices, the maximum number of edges is $n(n-1)$
 - example: G_1 is a complete graph

Adjacent and Incident

- If (v_0, v_1) is an edge in an undirected graph,
 - v_0 and v_1 are **adjacent**
 - The edge (v_0, v_1) is incident on vertices v_0 and v_1
- If $\langle v_0, v_1 \rangle$ is an edge in a directed graph
 - v_0 is **adjacent to** v_1 , and v_1 is **adjacent from** v_0
 - The edge $\langle v_0, v_1 \rangle$ is incident on v_0 and v_1

***Figure 6.3:**Example of a graph with feedback loops and a
multigraph (p.260)




Subgraph and Path


- A **subgraph** of G is a graph G' such that $V(G')$ is a subset of $V(G)$ and $E(G')$ is a subset of $E(G)$
- A **path** from vertex v_p to vertex v_q in a graph G , is a sequence of vertices, $v_p, v_{i1}, v_{i2}, \dots, v_{in}, v_q$, such that $(v_p, v_{i1}), (v_{i1}, v_{i2}), \dots, (v_{in}, v_q)$ are edges in an undirected graph
- The **length of a path** is the number of edges on it

Figure 6.4: subgraphs of G_1 and G_3 (p.261)


Go, change the world




G_1




(i)



(ii)



(iii)




(iv)

(a) Some of the subgraph of G_1




G_3




單一


(i)



(ii)



(iii)



分開


(iv)

(b) Some of the subgraph of G_3


Simple Path and Style

- A **simple path** is a path in which all vertices, except possibly the first and the last, are distinct
- A **cycle** is a simple path in which the first and the last vertices are the same
- In an undirected graph G , two **vertices**, v_0 and v_1 are **connected** if there is a path in G from v_0 to v_1
- An undirected **graph** is **connected** if, for every pair of distinct vertices v_i, v_j , there is a path from v_i to v_j

connected



G_1



G_2

tree (acyclic graph)


Connected Component

- A **connected component** of an undirected graph is a maximal connected subgraph.
- A **tree** is a graph that is connected and acyclic.
- A directed graph is **strongly connected** if there is a directed path from v_i to v_j and also from v_j to v_i .
- A **strongly connected component** is a maximal subgraph that is strongly connected.

***Figure 6.5: A graph with two connected components (p.262)**

Go, change the world



connected component (maximal connected subgraph)



G_4 (not connected)

***Figure 6.6: Strongly connected components of G_3 (p.262)**

not strongly connected strongly connected component
(maximal strongly connected subgraph)



Degree



- The **degree** of a vertex is the number of edges incident to that vertex
- For directed graph,
 - the **in-degree** of a vertex v is the number of edges that have v as the head
 - the **out-degree** of a vertex v is the number of edges that have v as the tail
 - if d_i is the degree of a vertex i in a graph G with n vertices and e edges, the number of edges is

$$e = \left(\sum_{i=0}^{n-1} d_i \right) / 2$$


undirected graph

Go, change the world

degree



directed graph
in-degree
out-degree



G_3

in: 1, out: 1

in: 1, out: 2

in: 1, out: 0



Graph Representations


- Adjacency Matrix
- Adjacency Lists
- Adjacency Multilists

Adjacency Matrix

- Let $G=(V,E)$ be a graph with n vertices.
- The **adjacency matrix** of G is a two-dimensional n by n array, say `adj_mat`
- If the edge (v_i, v_j) is in $E(G)$, $\text{adj_mat}[i][j]=1$
- If there is no such edge in $E(G)$, $\text{adj_mat}[i][j]=0$
- The adjacency matrix for an undirected graph is symmetric; the adjacency matrix for a digraph need not be symmetric

Examples for Adjacency Matrix

Go, change the world



$$\begin{bmatrix} 0 & 1 & 1 & 1 \\ 1 & 0 & 1 & 1 \\ 1 & 1 & 0 & 1 \\ 1 & 1 & 1 & 0 \end{bmatrix}$$



G_1



$$\begin{bmatrix} 0 & 1 & 0 \\ 1 & 0 & 1 \\ 0 & 0 & 0 \end{bmatrix}$$

G_2

symmetric



$$\begin{bmatrix} 0 & 1 & 1 & 0 & 0 & 0 & 0 & 0 \\ 1 & 0 & 0 & 1 & 0 & 0 & 0 & 0 \\ 1 & 0 & 0 & 1 & 0 & 0 & 0 & 0 \\ 0 & 1 & 1 & 0 & 0 & 0 & 0 & 0 \\ 0 & 0 & 0 & 0 & 0 & 1 & 0 & 0 \\ 0 & 0 & 0 & 0 & 1 & 0 & 1 & 0 \\ 0 & 0 & 0 & 0 & 0 & 1 & 0 & 1 \\ 0 & 0 & 0 & 0 & 0 & 0 & 1 & 0 \end{bmatrix}$$

G_4

undirected: $n^2/2$

directed: n^2

Merits of Adjacency Matrix



- From the adjacency matrix, to determine the connection of vertices is easy
- The degree of a vertex is $\sum_{j=0}^{n-1} adj_m$
- For a digraph, the row sum is the out_degree, while the column sum is the in_degree

$$ind(vi) = \sum_{j=0}^{n-1} A[j, i] \quad outd(vi) = \sum_{j=0}^{n-1} A[i, j]$$


Data Structures for Adjacency Lists

Each row in adjacency matrix is represented as an adjacency list.



```
#define MAX_VERTICES 50
typedef struct node *node_pointer;
typedef struct node {
    int vertex;
    struct node *link;
};
node_pointer graph[MAX_VERTICES];
int n=0; /* vertices currently in use *
```




G_1



G_3



Go, change the world



G_4

An undirected graph with n vertices and e edges \Rightarrow n head nodes and $2e$ list nodes

Interesting Operations

- **degree of a vertex** in an undirected graph

 - # of nodes in adjacency list

- **# of edges** in a graph

 - determined in $O(n+e)$

- **out-degree** of a vertex in a directed graph


 - # of nodes in its adjacency list

- **in-degree** of a vertex in a directed graph

 - traverse the whole data structure

Compact Representation


Go, change the world



node[0] ... node[n-1]: starting point for vertices
node[n]: n+2e+1
node[n+1] ... node[n+2e]: head node of edge

[0]	9		[8]	23		[16]	2
[1]	11	0	[9]	1	4	[17]	5
[2]	13		[10]	2	5	[18]	4
[3]	15	1	[11]	0		[19]	6
[4]	17		[12]	3	6	[20]	5
[5]	18	2	[13]	0		[21]	7
[6]	20		[14]	3	7	[22]	6
[7]	22	3	[15]	1			

Figure 6.10: Inverse adjacency list for G_3



Determine in-degree of a vertex in a fast way.

Figure 6.11: Alternate node structure for adjacency lists (p.267)

tail	head	column link for head	row link for tail
------	------	----------------------	-------------------

Figure 6.12: Orthogonal representation for graph G_3 (p.268)

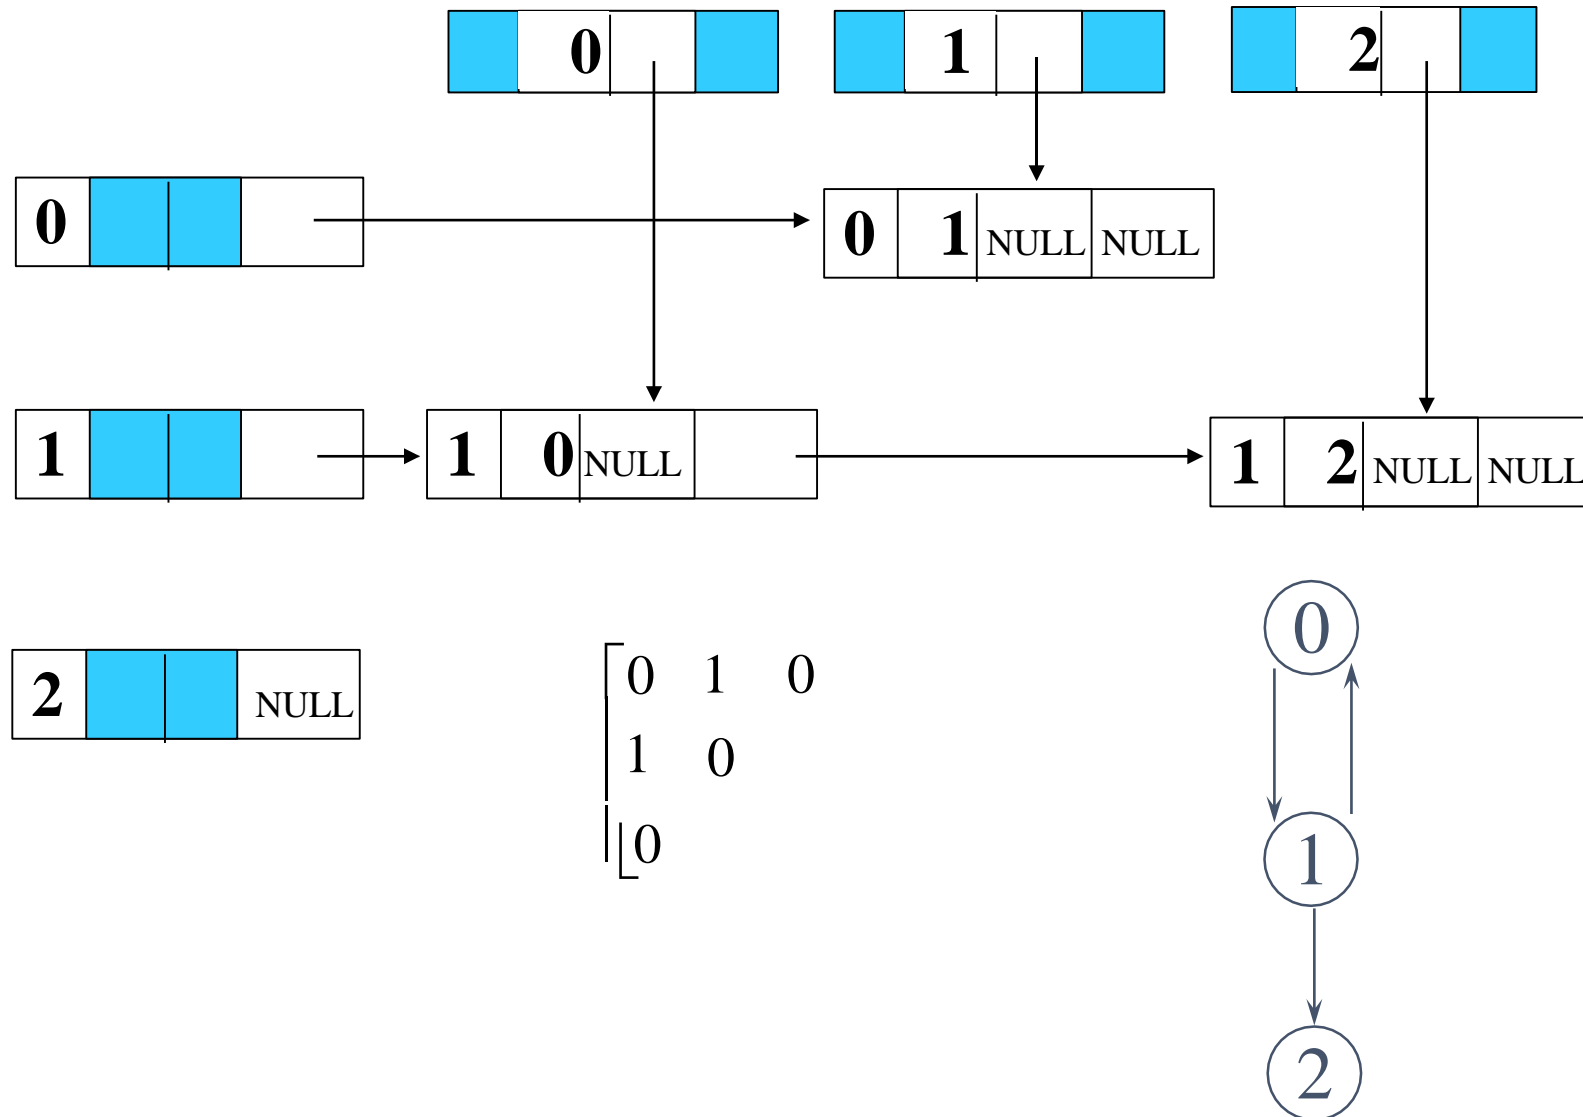




Figure 6.13: Alternate order adjacency list for G_1 (p.268)

Order is of no significance.



□ An edge in an undirected graph is represented by two nodes in adjacency list representation.

□ Adjacency Multilists

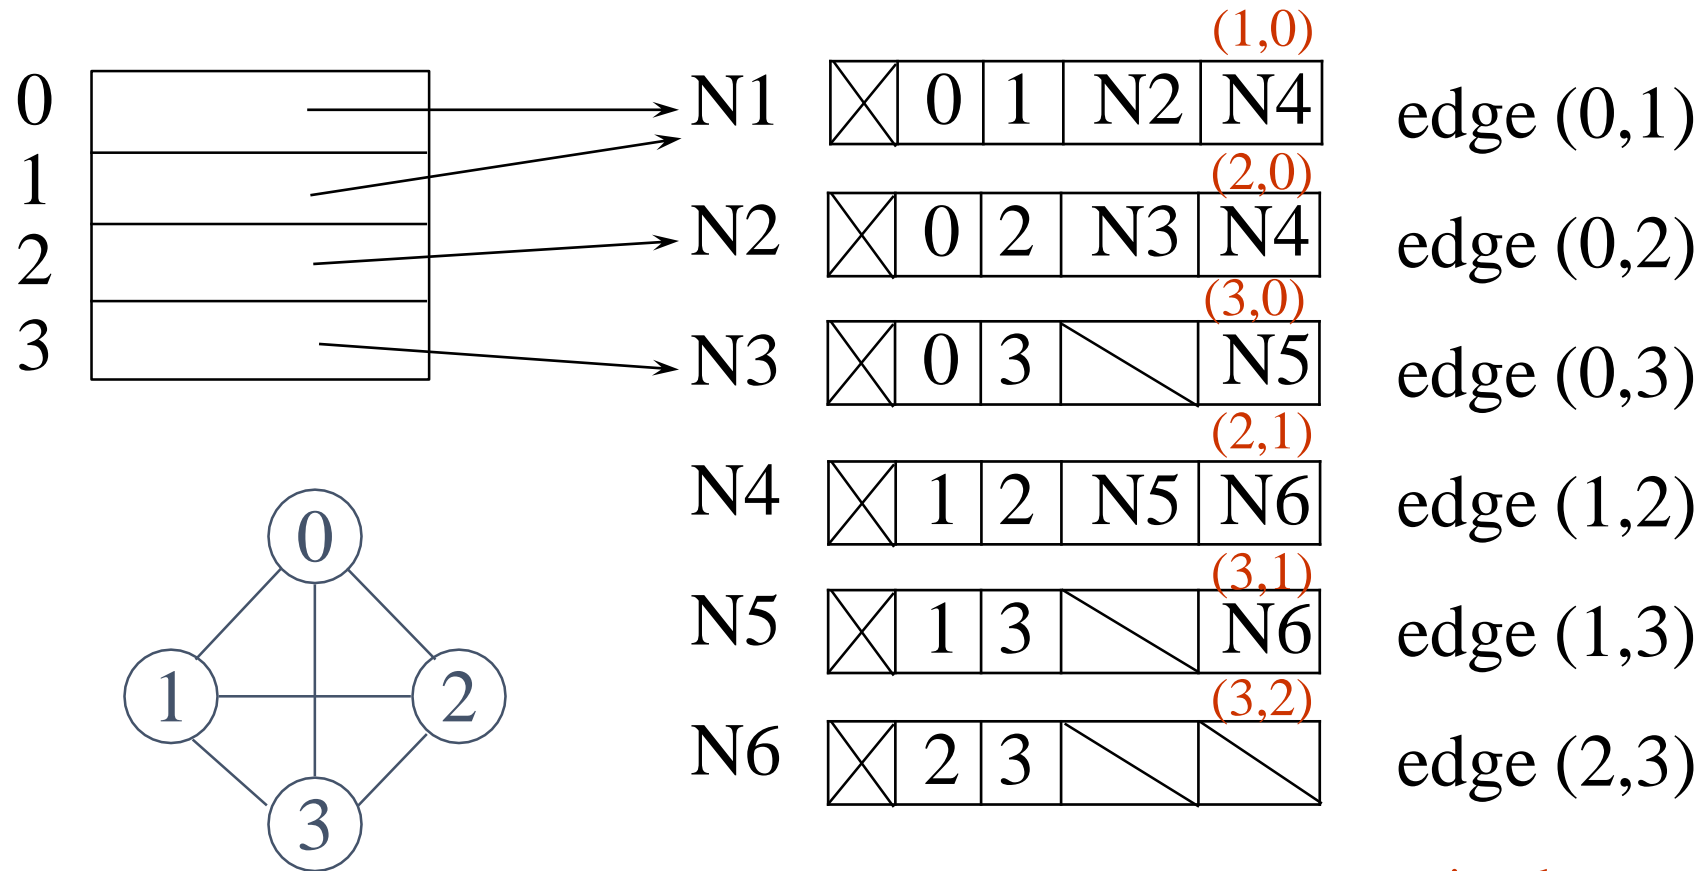
–lists in which nodes may be shared among several lists.

(an edge is shared by two different paths)

marked	vertex1	vertex2	path1	path2
--------	---------	---------	-------	-------

Example for Adjacency Multilists

Lists: vertex 0: M1->M2->M3, vertex 1: M1->M4->M5
vertex 2: M2->M4->M6, vertex 3: M3->M5->M6



six edges

Adjacency Multilists

```
typedef struct edge *edge_pointer;  
typedef struct edge {  
    short int marked;  
    int vertex1, vertex2;  
    edge_pointer path1, path2;  
};  
edge_pointer graph[MAX_VERTICES];
```

marked	vertex1	vertex2	path1	path2
--------	---------	---------	-------	-------

Some Graph Operations

- Traversal

Given $G=(V,E)$ and vertex v , find all $w \in V$, such that w connects v .

- Depth First Search (DFS)

- preorder tree traversal

- Breadth First Search (BFS)

- level order tree traversal

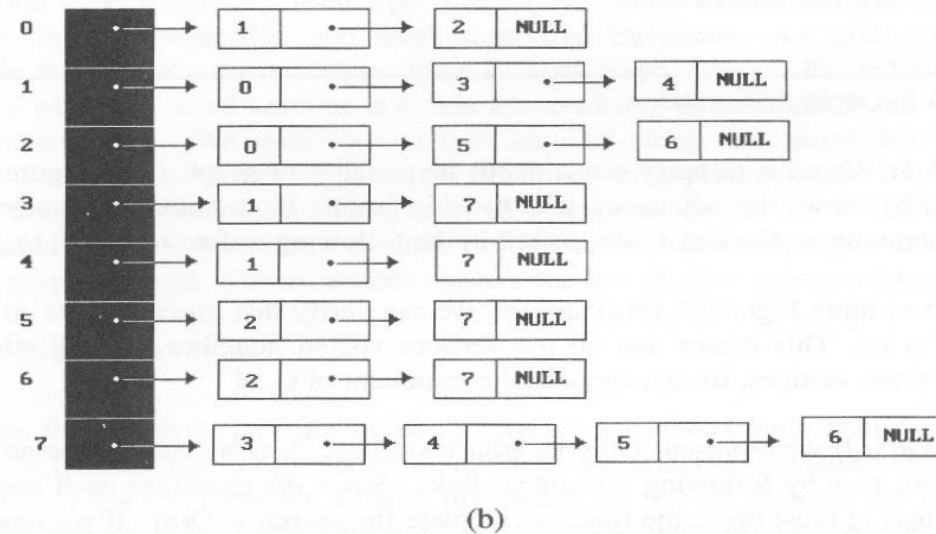

- Connected Components

- Spanning Trees

***Figure 6.19: Graph G and its adjacency lists (p.274)**

Go, change the world

depth first search: $v_0, v_1, v_3, v_7, v_4, v_5, v_2, v_6$



breadth first search: $v_0, v_1, v_2, v_3, v_4, v_5, v_6, v_7$

Depth First Search

Go, change the world

```
#define FALSE 0
#define TRUE 1
short int visited[MAX_VERTICES];
```

```
void dfs(int v)
{
    node_pointer w;
    visited[v]= TRUE;
    printf("%5d", v);
    for (w=graph[v]; w; w=w->link)
        if (!visited[w->vertex])
            dfs(w->vertex);
}
```

Data structure

adjacency list: $O(e)$

adjacency matrix: $O(n^2)$

Breadth First Search

```
typedef struct queue *queue_pointer;  
typedef struct queue {  
    int vertex;  
    queue_pointer link;  
};  
void addq(queue_pointer *,  
          queue_pointer *, int);  
int deleteq(queue_pointer *);
```

Breadth First Search *(Continued)*

```
void bfs(int v)
{
    node_pointer w;
    queue_pointer front, rear;
    front = rear = NULL;
    printf("%5d", v);
    visited[v] = TRUE;
    addq(&front, &rear, v);
```

adjacency list: $O(e)$

adjacency matrix: $O(n^2)$

```
while (front) {  
    v= deleteq(&front);  
    for (w=graph[v]; w; w=w->link)  
        if (!visited[w->vertex]) {  
            printf("%5d", w->vertex);  
            addq(&front, &rear, w->vertex);  
            visited[w->vertex] = TRUE;  
        }  
    }  
}
```

Connected Components


```
void connected(void)
{
    for (i=0; i<n; i++) {
        if (!visited[i]) {
            dfs(i);
            printf("\n");
        }
    }
}
```

adjacency list: $O(n+e)$
adjacency matrix: $O(n^2)$




Spanning Trees

- When graph G is connected, a depth first or breadth first search starting at any vertex will visit all vertices in G
- A **spanning tree** is any tree that consists solely of edges in G and that includes all the vertices
- $E(G): T$ (tree edges) + N (nontree edges)
where T : set of edges used during search
 N : set of remaining edges

Examples of Spanning Tree



G_1




Possible spanning trees

Spanning Trees

- Either dfs or bfs can be used to create a spanning tree
 - When dfs is used, the resulting spanning tree is known as a **depth first spanning tree**
 - When bfs is used, the resulting spanning tree is known as a **breadth first spanning tree**
- While adding a nontree edge into any spanning tree, this will create a cycle

DFS VS BFS Spanning Tree





A spanning tree is a **minimal subgraph**, G' , of G such that $V(G')=V(G)$ and G' is connected.

Any connected graph with **n** vertices must have at least **$n-1$** edges.

A **biconnected graph** is a connected graph that has no **articulation points**.


biconnected graph






connected graph


two connected components



one connected graph



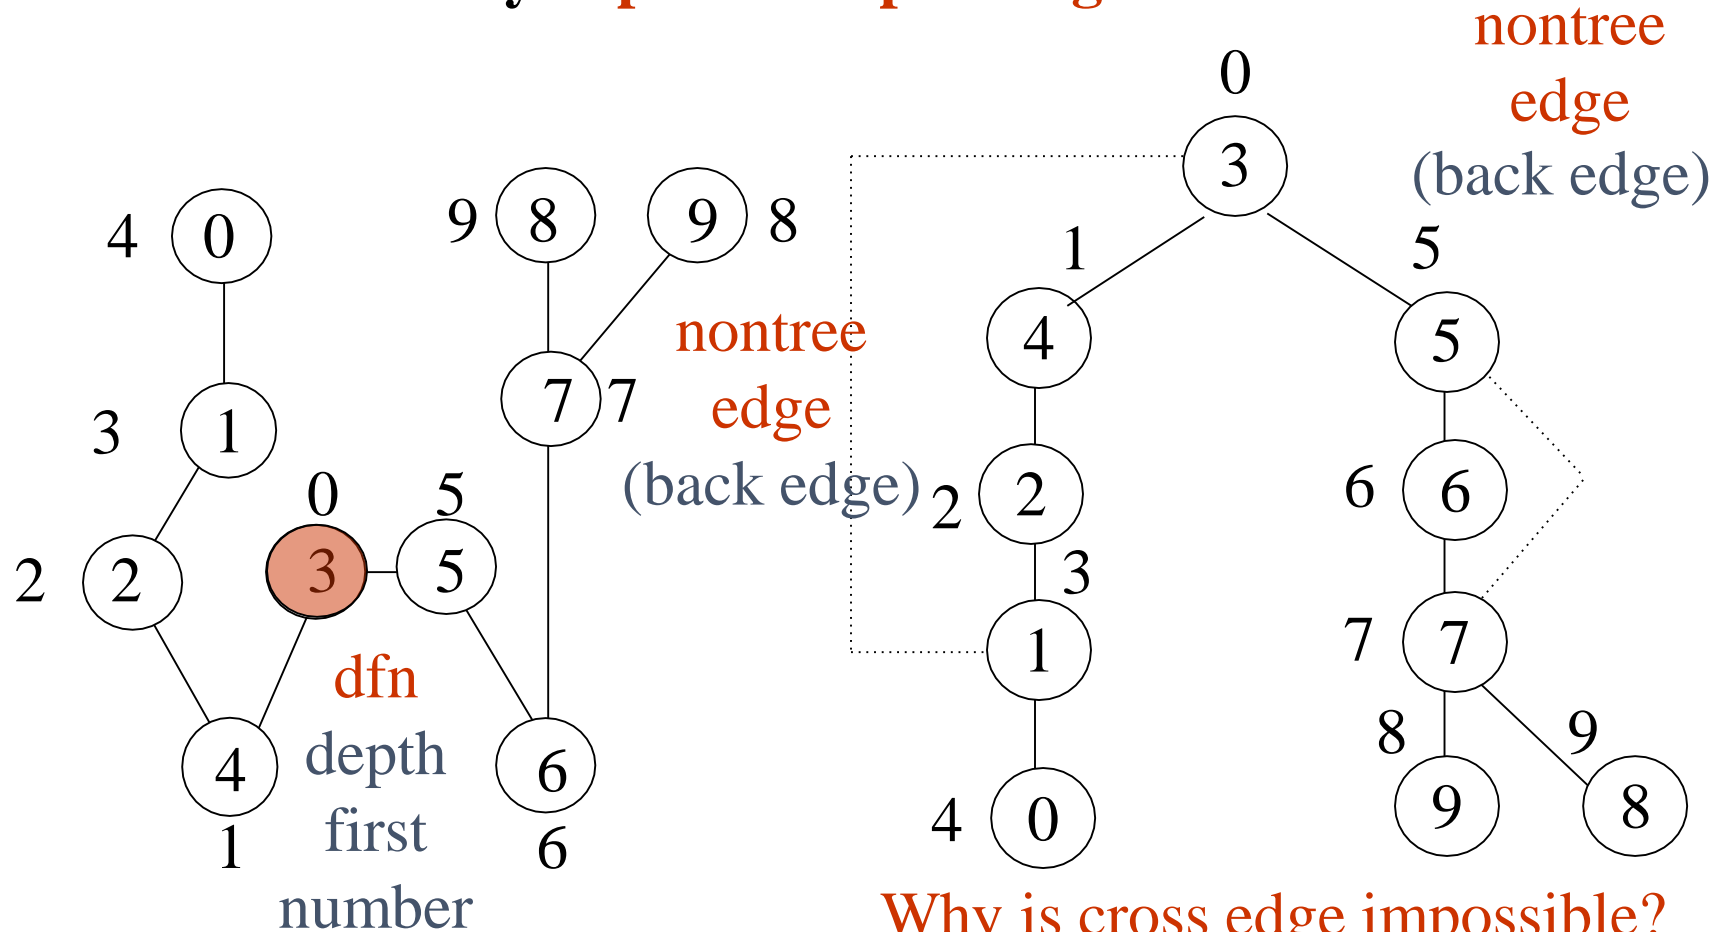
biconnected component: a maximal connected subgraph H
(no subgraph that is both biconnected and properly contains H)



biconnected components

Find biconnected component of a connected undirected graph by **depth first spanning tree**

Go, change the world




(a) depth first spanning tree

(b)

If u is an ancestor of v then $\text{dfn}(u) < \text{dfn}(v)$.

***Figure 6.24:** *dfn* and *low* values for *dfs* spanning tree with *root* = 3(p.281)

Vertex	0	1	2	3	4	5	6	7	8	9
<i>dfn</i>	4	3	2	0	1	5	6	7	9	8
<i>low</i>	4	0	0	0	0	5	5	5	9	8



*The root of a depth first spanning tree is an articulation point iff it has at least two children.

*Any other vertex u is an articulation point iff it has at least one child w such that we cannot reach an ancestor of u using a path that consists of


(1) only w (2) descendants of w (3) single back edge.

$\text{low}(u) = \min\{\text{dfn}(u),$
 $\min\{\text{low}(w) \mid w \text{ is a child of } u\},$
 $\min\{\text{dfn}(w) \mid (u, w) \text{ is a back edge}\}$

u : articulation point

$\text{low}(\text{child}) \geq \text{dfn}(u)$

vertex	dfn	low	child	low_child	low:dfn
0	4	4 (4,n,n)	null	null	null:4
1	3	0 (3,4,0)	0	4	4 ≥ 3 •
2	2	0 (2,0,n)	1	0	0 < 2
3	0	0 (0,0,n)	4,5	0,5	0,5 ≥ 0 •
4	1	0 (1,0,n)	2	0	0 < 1
5	5	5 (5,5,n)	6	5	5 ≥ 5 •
6	6	5 (6,5,n)	7	5	5 < 6
7	7	5 (7,8,5)	8,9	9,8	9,8 ≥ 7 •
8	9	9 (9,n,n)	null	null	null, 9
9	8	8 (8,n,n)	null	null	null, 8





*Program 6.5: Initializaiton of *dfn* and *low* (p.282)

```
void init(void)
{
    int i;
    for (i = 0; i < n; i++) {
        visited[i] = FALSE;
        dfn[i] = low[i] = -1;
    }
    num = 0;
}
```

```

void dfnlow(int u, int v)                                Initial call: dfn(x,-1)
{
  /* compute dfn and low while performing a dfs search
     beginning at vertex u, v is the parent of u (if any) */
  node_pointer ptr;
  int w;
  dfn[u] = low[u] = num++;                               low[u]=min{ dfn(u), ...}
  for (ptr = graph[u]; ptr; ptr = ptr->link) {
    w = ptr->vertex;
    if (dfn[w] < 0) { /*w is an unvisited vertex */
      dfnlow(w, u);
      low[u] = MIN2(low[u], low[w]);
    }
    low[u]=min{..., min{low(w)|w is a child of u}, ...}
    else if (w != v) dfn[w]≠0 非第一次, 表示藉back edge
  }
  low[u]=min{..., ..., min{dfn(w)|(u,w) is a back edge}
}

```

O X



```
void bicon(int u, int v)
{
    /* compute dfn and low, and output the edges of G by their
       biconnected components, v is the parent (if any) of the u
       (if any) in the resulting spanning tree. It is assumed that all
       entries of dfn[ ] have been initialized to -1, num has been
       initialized to 0, and the stack has been set to empty */
    node_pointer ptr;
    int w, x, y;
    dfn[u] = low[u] = num++;    low[u]=min{ dfn(u), ...}
    for (ptr = graph[u]; ptr; ptr = ptr->link) {
        w = ptr->vertex;
        if ( v != w && dfn[w] < dfn[u] ) (1) dfn[w]=-1 第一次
        (2) dfn[w]!=-1非第一次, 藉back
        add(&top, u, w);    /* add edge to stack */    edge
    }
```

```
if(dfn[w] < 0) {
    bicon(w, u);
    low[u] = MIN2(low[u], low[w]);
    if (low[w] >= dfn[u] ){ articulation point
        printf("New biconnected component: ");
        do { /* delete edge from stack */
            delete(&top, &x, &y);
            printf(" <%d, %d>" , x, y);
        } while (!(( x == u) && (y == w)));
        printf("\n");
    }
}
else if (w != v) low[u] = MIN2(low[u], dfn[w]);
}
}

low[u]=min{..., ..., min{dfn(w)|(u,w) is a back edge} }
```

Minimum Cost Spanning Tree

- The cost of a spanning tree of a weighted undirected graph is the sum of the costs of the edges in the spanning tree
- A minimum cost spanning tree is a spanning tree of least cost
- Three different algorithms can be used
 - Kruskal
 - Prim
 - Sollin

Select $n-1$ edges from a weighted graph of n vertices with minimum cost.

Greedy Strategy

- An optimal solution is constructed in stages
- At each stage, the best decision is made at this time
- Since this decision cannot be changed later, we make sure that the decision will result in a feasible solution
- Typically, the selection of an item at each stage is based on a least cost or a highest profit criterion

Kruskal's Idea

Build a minimum cost spanning tree T by adding edges to T one at a time

Select the edges for inclusion in T in nondecreasing order of the cost

An edge is added to T if it does not form a cycle

Since G is connected and has $n > 0$ vertices, exactly $n-1$ edges will be selected

Examples for Kruskal's Algorithm

Go, change the world

~~0-10-5~~

~~2-12-3~~

~~1-14-6~~

~~1-16-2~~



~~3-18-6~~

~~3-22-4~~

~~4-24-6~~

~~4-25-5~~

~~0-28-1~~



0-10-5

2-12-3

1-14-6

1-16-2




3-18-6

3-22-4

4-24-6

4-25-5

0-28-1



↓ + 3-6
cycle

0-10-5

2-12-3

1-14-6

1-16-2


3-18-6

3-22-4

4-24-6

4-25-5

0-28-1




+

4-6

cycle

cost = 10 + 25 + 22 + 12 + 16 + 14



Kruskal's Algorithm

```

T= {};
while (T contains less than n-1 edges
      && E is not empty) {
  choose a least cost edge (v,w) from E;
  delete (v,w) from E; min heap construction time  $O(e)$ 
  if ((v,w) does not create a cycle in T) choose and delete  $O(\log e)$ 
    add (v,w) to T
  else discard (v,w); find find & union  $O(\log e)$ 
} {0,5}, {1,2,3,6}, {4} + edge(3,6) X + edge(3,4) --> {0,5}, {1,2,3,4,6}
if (T contains fewer than n-1 edges)
  printf("No spanning tree\n");
 $O(e \log e)$ 

```

Prim's Algorithm


Go, change the world





(tree all the time vs. forest)

```
T = { } ;  
TV = { 0 } ;  
while (T contains fewer than n-1 edges)  
{  
    let (u,v) be a least cost edge such that  $u \in$  and  
         $v \notin$   
    if (there is no such edge ) break; add v to TV;  
    add (u,v) to T;  
}  
if (T contains fewer than n-1 edges)  
    printf("No spanning tree\n");
```

Examples for Prim's Algorithm


Go, change the world






Sollin's Algorithm

Go, change the world




vertex	edge
0	0 -- 10 --> 5, 0 -- 28 --> 1
1	1 -- 14 --> 6, 1 -- 16 --> 2, 1 -- 28 --> 0
2	2 -- 12 --> 3, 2 -- 16 --> 1
3	3 -- 12 --> 2, 3 -- 18 --> 6, 3 -- 22 --> 4
4	4 -- 22 --> 3, 4 -- 24 --> 6, 5 -- 25 --> 5
5	5 -- 10 --> 0, 5 -- 25 --> 4
6	6 -- 14 --> 1, 6 -- 18 --> 3, 6 -- 24 --> 4



Single Source All Destinations

Determine the shortest paths from v_0 to all the remaining vertices.

*Figure 6.29: Graph and shortest paths from v_0 (p.293)



(a)

path	length
1) $v_0 v_2$	10
2) $v_0 v_2 v_3$	25
3) $v_0 v_2 v_3 v_1$	45
4) $v_0 v_4$	45

(b)

All Pairs Shortest Paths

Find the shortest paths between all pairs of vertices.

Solution 1

- Apply **shortest path** n times with each vertex as source.

$O(n^3)$

Solution 2


- Represent the graph G by its cost adjacency matrix with $\text{cost}[i][j]$
- If the edge $\langle i, j \rangle$ is not in G , the $\text{cost}[i][j]$ is set to some sufficiently large number
- $A[i][j]$ is the cost of the shortest path from i to j , using only those intermediate vertices with an index $\leq k$

All Pairs Shortest Paths (*Continued*)

- The cost of the shortest path from i to j is $A^{n-1}[i][j]$, as no vertex in G has an index greater than $n-1$
- $A^{-1}[i][j] = \text{cost}[i][j]$
- Calculate the $A^0, A^1, A^2, \dots, A^{n-1}$ from A^{-1} iteratively
- $A^k[i][j] = \min\{A^{k-1}[i][j], A^{k-1}[i][k] + A^{k-1}[k][j]\}, k \geq 0$

Graph with negative cycle

Go, change the world



(a) Directed graph


$$\begin{bmatrix} 0 & 1 & \infty \\ -2 & 0 & 1 \\ \infty & \infty & 0 \end{bmatrix}$$

(b) A^{-1}

The length of the shortest path from vertex 0 to vertex 2 is $-\infty$.

$0, 1, 0, 1, 0, 1, \dots, 0, 1, 2$

* Figure 6.33: Directed graph and its cost matrix (p.299)




(a) Digraph G

	0	1	2
0	0	4	11
1	6	0	2
2	3	∞	0

(b) Cost adjacency matrix for G

$$A^{-1}$$


	0	1	2
0	0	4	11
1	6	0	2
2	3	∞	0



$$A^0$$


	0	1	2
0	0	4	1
1	6	0	
2	3		

Go, change the world




$$A^1$$

	0	1	2
0	0	4	6
1	6	0	2
2	3	7	0



$$A^2$$

	0	1	2
0	0	4	6
1	5	0	2
2	3	7	0



Transitive Closure

Go, change the world


Goal: given a graph with unweighted edges, determine if there is a path from i to j for all i and j .

(1) Require positive path (> 0) lengths.

transitive closure matrix

(2) Require nonnegative path (≥ 0) lengths.

reflexive transitive closure matrix



(a) Digraph G

$$\begin{matrix} 0 \\ 1 \\ 2 \\ 3 \\ 4 \end{matrix} \begin{bmatrix} 0 & 1 & 0 & 0 & 0 \\ 0 & 0 & 1 & 0 & 0 \\ 0 & 0 & 0 & 1 & 0 \\ 0 & 0 & 0 & 0 & 1 \\ 0 & 0 & 1 & 0 & 0 \end{bmatrix}$$

(b) Adjacency matrix A for G

$$\begin{matrix} 0 \\ 1 \\ 2 \\ 3 \\ 4 \end{matrix} \begin{bmatrix} 0 & 1 & 1 & 1 & 1 \\ 0 & 0 & 1 & 1 & 1 \\ 0 & 0 & 1 & 1 & 1 \\ 0 & 0 & 1 & 1 & 1 \\ 0 & 0 & 1 & 1 & 1 \end{bmatrix}$$

cycle

(c) transitive closure matrix A^+

There is a path of length > 0

$$\begin{matrix} 0 \\ 1 \\ 2 \\ 3 \\ 4 \end{matrix} \begin{bmatrix} 1 & 1 & 1 & 1 & 1 \\ 0 & 1 & 1 & 1 & 1 \\ 0 & 0 & 1 & 1 & 1 \\ 0 & 0 & 1 & 1 & 1 \\ 0 & 0 & 1 & 1 & 1 \end{bmatrix}$$

reflexive

(d) reflexive transitive closure matrix A^*

There is a path of length ≥ 0